Performance Metrics:
** Latencies are in msec.

		Server Threads			
Client Threads|	Info	50	100	150	200
10	Average	32	34	19	24
	Median	2	2	2	3
	90%	4	4	5	8
	95%	6	5	7	12
	99%	36	12	24	40
	Through Put /sec	289.9	268.5	474.0	369.1
	Error %	0.0%	0.0%	0.0%	0.0%
20	Average	65	39	39	38
	Median	6	5	4	5
	90%	12	9	7	10
	95%	17	12	10	14
	99%	42	33	40	39
	Through Put /sec	379.1	475.0	467.3	476.5
	Error %	0.00%	0.0%	0.0%	0.0%
40	Average	77	127	139	142
	Median	9	6	7	7
	90%	14	12	12	15
	95%	17	23	27	27
	99%	61	76	2017	2233
	Through Put /sec	477.6	287.9	265.4	259.6

	Error %	0.24%	0.23%	0.23%	0.23%

For the current system configuration, Thresholding on the 99% latency. It seems like crossing after 150 level, resulting in increase of the 99%. Increasing threads beyond is not resulting in performance improvement. These is because of the latencies due to context switch between the threads. Increasing thread count has given us better performance till a limit. But later the performance is not improving.

Increasing to 200 has not resulted in a drastic improvement. Its similar to 150, but with higher latencies as well. So For the current system configuration on which this server is tested, 150 threads would be a proper choice.

System Configuration:

MacBook Pro (13-inch, 2017, Two Thunderbolt 3 ports)

Memory : 8 GB 2133 MHz LPDDR3

Processor: 2.3 GHz Intel Core i5